

Particle Filters

10 April 2018

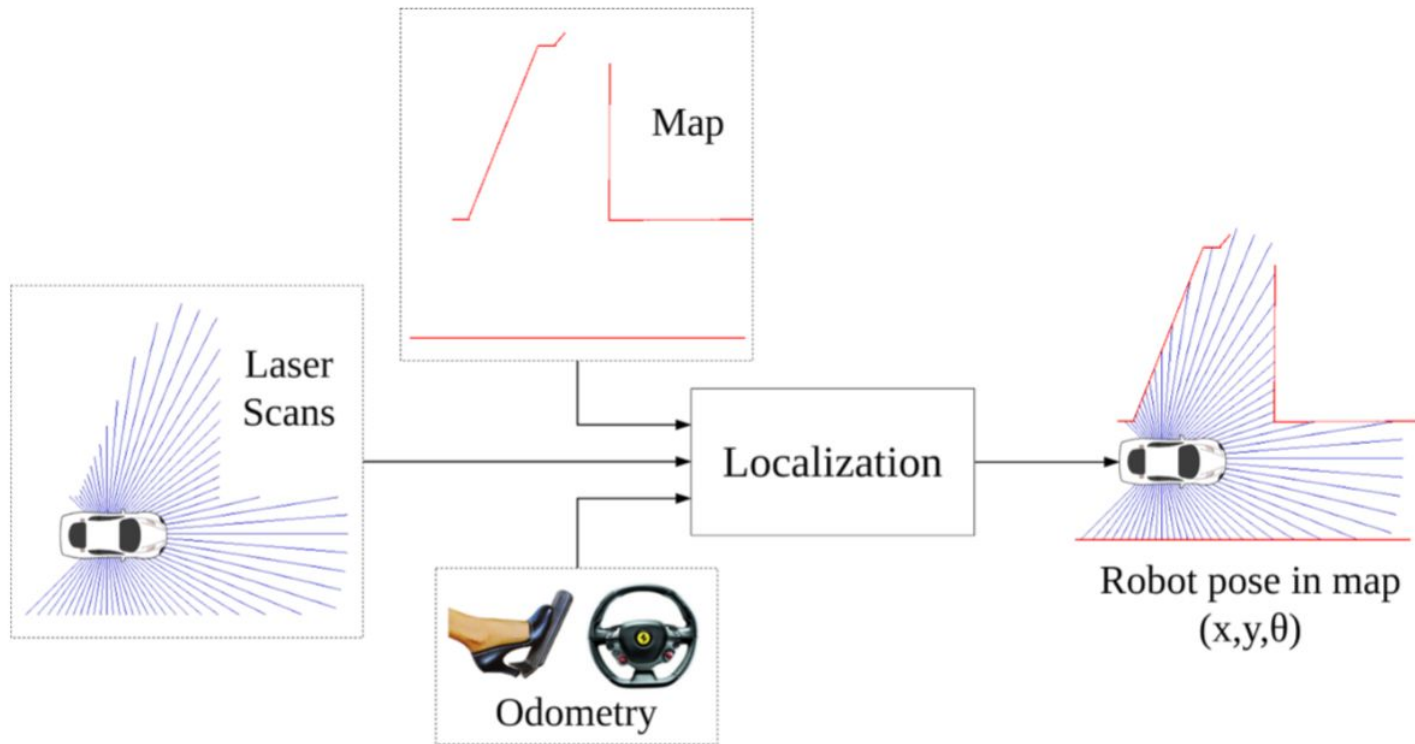
Outline

- What is it?
- Real world examples
- Key Concepts
- The PF Algorithm
- PF Algorithm Review Videos
- Resampling Algorithms
- Compare PF with other Probabilistic Filters

Outline

- **What is it?**
- Real world examples
- Key Concepts
- The PF Algorithm
- PF Algorithm Review
- Resampling Algorithms
- Compare PF with other Probabilistic Filters

What is it?



What is it?

A probabilistic algorithm that uses

- **"particles"**
- **"importance/weights"**, *and*
- **"resampling"**

to estimate the current state/location of a moving object.

Outline

- What is it?
- **Real world examples**
- Key Concepts
- The PF Algorithm
- PF Algorithm Review
- Resampling Algorithms
- Compare PF with other Probabilistic Filters

Real World Examples



Figure 2: Particle filters have been used successfully for on-board localization of soccer-playing Aibo robots with as few as 50 particles [26].

Real World Examples




Figure 7: Particle filter-based people tracker: This algorithm uses maps to simultaneously localize a moving robot and an unknown number of nearby people. This sequence shows the evolution of the conditional particle filter from global uncertainty to successful localization and tracking of the robot.

Real World Examples

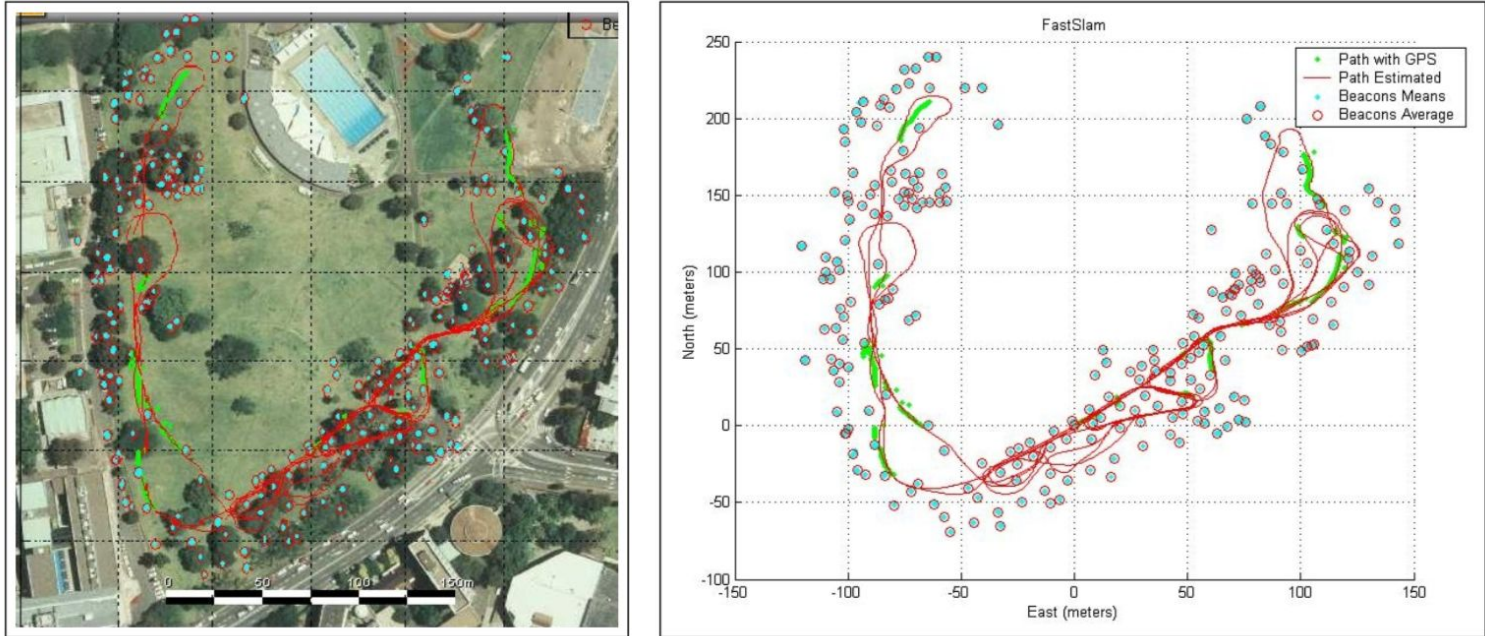


Figure 5: FastSLAM with real-world data: Shown here is a map of an outdoor environment (Victoria Park in Sydney), along with GPS information displayed here only for evaluating the accuracy of the resulting map. The resulting map error is extremely small, comparable in magnitude to the EKF solution. These results were obtained Juan Nieto, Eduardo Nebot, and Jose Guivant from the Australian Center of Field Robotics in Sydney, and are reprinted here with permission.

Real World Examples

Autonomous RC-Car MIT Video

Autonomous RC-Car UPenn Video


Outline

- What is it?
- Real world examples
- **Key Concepts**
- The PF Algorithm
- PF Algorithm Review
- Resampling Algorithms
- Compare PF with other Probabilistic Filters


Key Concepts (1 / 4)

- Each particle is a possible state of a system






Key Concepts (1 / 4)



Key Concepts (1 / 4)




Key Concepts (2/4)

- A particle becomes "more important" if its predicted observation matches the actual observation
- “Data Association”




x_{Particle}
 y_{Particle}
 $\ominus_{\text{Particle}}$
Weight


(a)



(b)



Importance Weight




Key Concepts (2/4)

- **Many ways to compute importance!**
- Univariate
- Bivariate
- Correlation

Key Concepts (2/4)

- Univariate




Key Concepts (2/4)

- Univariate

```
def gaussian_prob(mu, sigma, x):  
    # calculate the probability of x for 1-dim Gaussian with mean mu and var sigma  
    return exp( -((mu - x) ** 2) / (sigma ** 2) / 2.) / sqrt( 2. * pi * (sigma ** 2))  
  
def get_weight(my_measurements, ground_measurements, noise):  
  
    w = 1.  
    for my_distance, ground_distance in zip(my_measurements, ground_measurements):  
        w *= gaussian_prob(mu = my_distance, sigma = noise, x = ground_distance)  
  
    return w + 1.e-300 # avoid round-off to zero
```




Key Concepts (2/4)

- Univariate

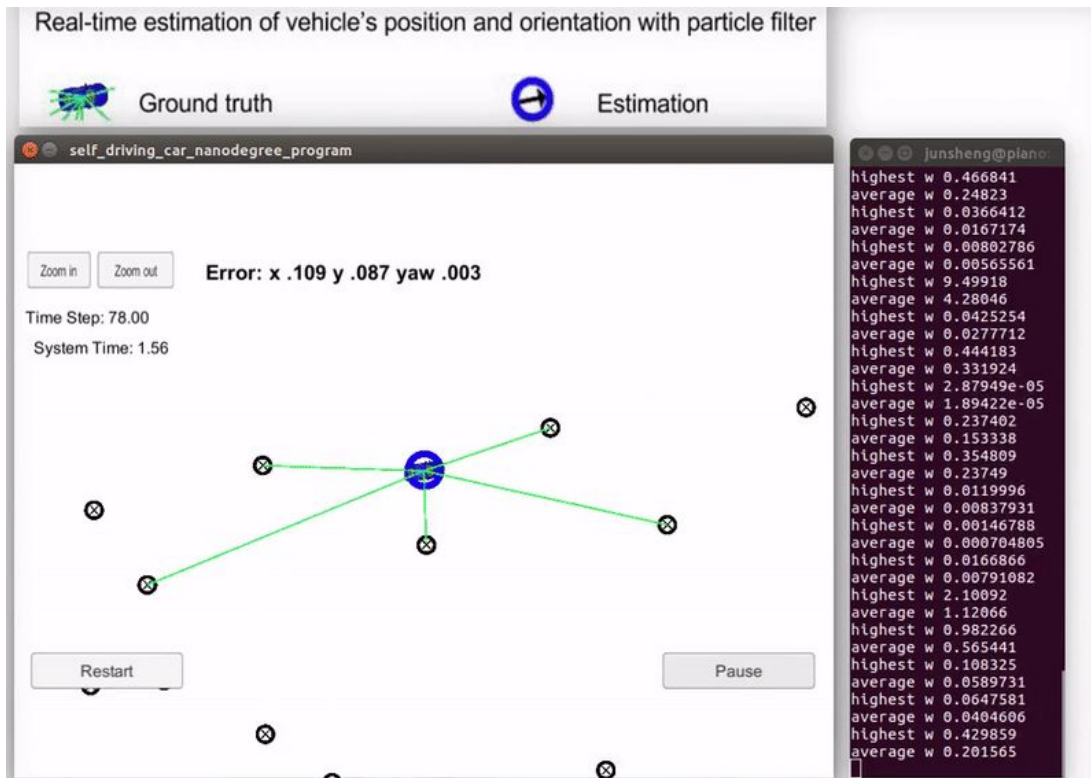


Key Concepts (2/4)

- Bivariate



Key Concepts (2/4)



Key Concepts (2/4)


- Bivariate

```
const double na = 0.5 / (stdx * stdx);  
const double nb = 0.5 / (stdy * stdy);  
const double d = sqrt( 2.0 * M_PI * stdx * stdy);  
  
246     const double dx = ox - predicted_x;  
247     const double dy = oy - predicted_y;  
248  
249     const double a = na * dx * dx;  
250     const double b = nb * dy * dy;  
251     const double r = exp(-(a + b)) / d;  
252     w *= r;
```

Key Concepts (2/4)


Scan Correlation

$$S = \frac{\sum_m \sum_n (A_{mn} - \bar{A})(B_{mn} - \bar{B})}{\sqrt{\left(\sum_m \sum_n (A_{mn} - \bar{A})^2\right) \left(\sum_m \sum_n (B_{mn} - \bar{B})^2\right)}}$$



Key Concepts (2/4)

- Scan Correlation




Key Concepts (3/4)

- Particles that are less important are more likely to die off
- Particles that are more important are more likely to clone itself

Key Concepts (3/4)

Replacement



N-Particles



RESAMPLING
N TIMES





Key Concepts (4/4)

- If our filter is working properly, as time goes by, ideally we become more and more certain



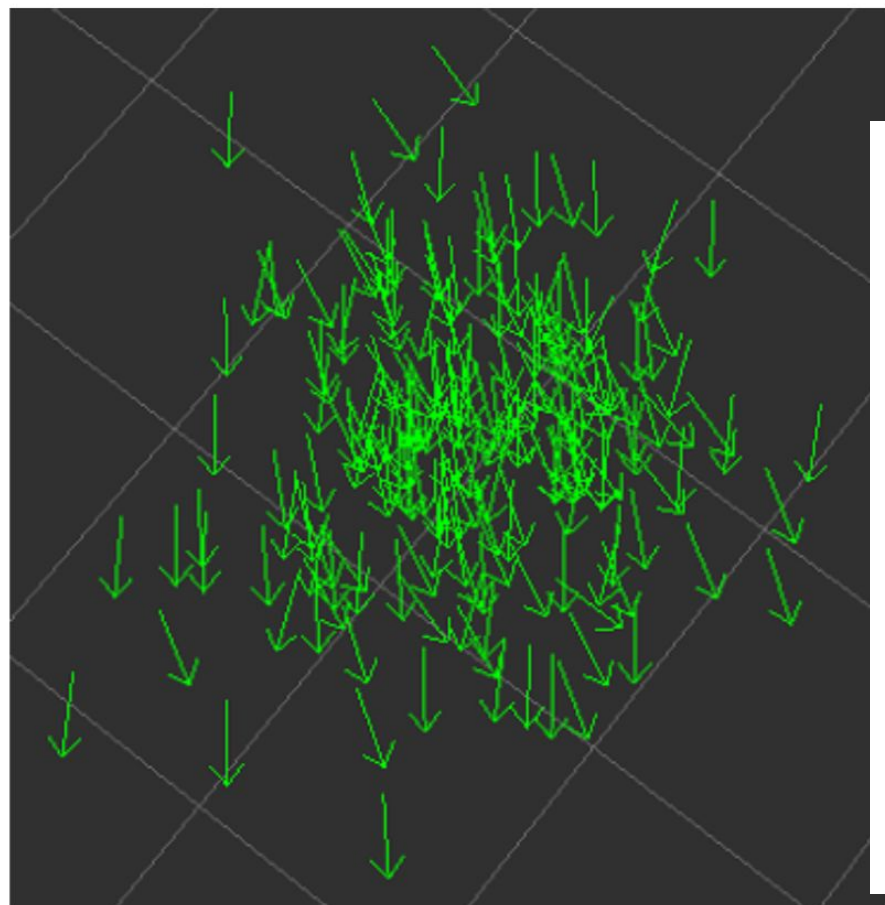


Fig. 3. Initial particle cloud generated based on a parametrized number of particles in the range between 25 and 200.

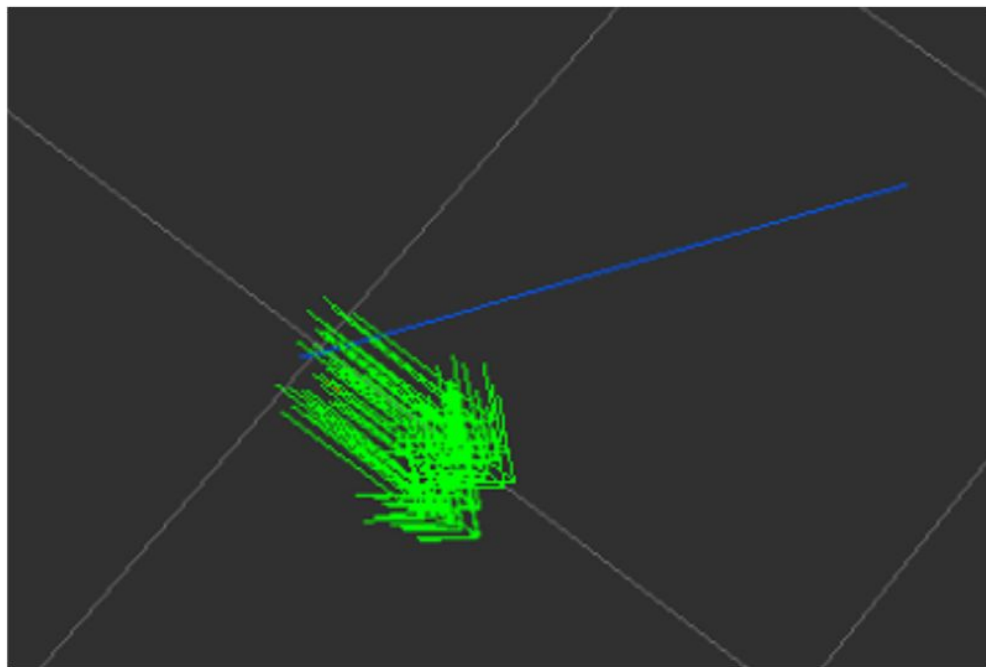
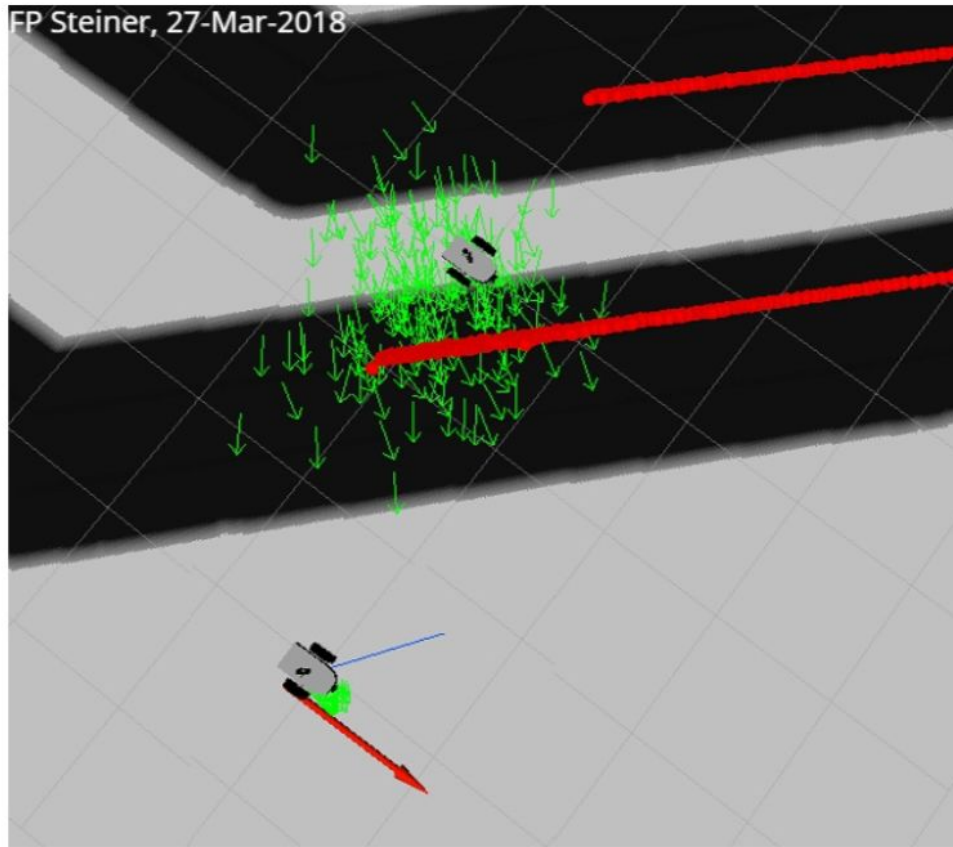


Fig. 4. Final particle cloud after the robot has reached its goal pose.

FP Steiner, 27-Mar-2018



FP Steiner, 27-Mar-2018




Fig. 9. Superposition of a screenshot of the initial and final pose and particle cloud of the classroom robot.

Fig. 10. Superposition of a screenshot of the initial and final pose and particle cloud of the redesigned robot.

Outline

- What is it?
- Real world examples
- Key Concepts
- **The PF Algorithm**
- The PF Algorithm Review videos
- Resampling Algorithms
- Compare PF with other Probabilistic Filters

Particle Filter Algorithm

STEPS

CONCEPTS

First Step

Previous Belief

Second Step

Motion Update

Third Step

Measurement Update

Fourth Step

Resampling


Fifth Step


New Belief

Particle Filter Algorithm

Previous Belief

1. Start with the current set of N particles (equal importance)







Particle Filter Algorithm

Motion Update

2. Move each particles as your robot moved (add noise)





Particle Filter Algorithm

Measurement Update

3. "Data Association"

- Get actual observation from "noisy" sensors
- For each particle,
 - Compare the actual observation, and the particle's predicted observation **(include some noise)**
 - If the predicted observation of your particle matches your actual observation then that particle is more important
 - If not, then consequently, less important.



*A bunch of
unlikely particles*

*Some very likely
particles*

Now, let's do the resampling!!

Particle Filter Algorithm

Resampling

4. Resample (You can think of it this way)

- Imagine you put all your particles[1...N] in a bag
- The ones that are more important are more likely to get picked
- Each time you pick a particle, you "clone it".
- Put the clone on your "new set" of particles
- Put the particle back in the bag, and pick one again
- Do this N times until you get a new set of particles



*A bunch of
unlikely particles*

*Some very likely
particles*

Now, lets do the resampling!!



~~A bunch of
unlikely particles~~

No resampled
particles!

~~Some very likely
particles~~

A lot of new particles
(plotted on top of each
other, so you can't see
them all...)

Particle Filter Algorithm

New Belief

5. Discard your old set of particles, this new set of particles will be your current set.



~~A bunch of
unlikely particles~~

No resampled
particles!

~~Some very likely
particles~~

A lot of new particles
(plotted on top of each
other, so you can't see
them all...)

Outline



- What is it?
- Real world examples
- Key Concepts
- The PF Algorithm
- **PF Algorithm Review**
- Resampling Algorithms
- Compare PF with other Probabilistic Filters

Particle Filter Algorithm Review

- Video from Udacity
- My Video: Four Landmark Example (Noisy Sensor, Noisy Motion)

Kullback–Leibler divergence (KLD Sampling)

- Variable Particle size
- Sample size is proportional to error between odometry position and sample based approximation
- i.e smaller sample size when particles have converged



Outline

- What is it?
- Real world examples
- Key Concepts
- The PF Algorithm
- PF Algorithm Review
- **Resampling Algorithms**
- Compare PF with other Probabilistic Filters

Resampling Algorithms

- Affects the performance of the filter.
- Research on the topic continues

Resampling Algorithms

Properties.

- preferentially select particles that have a higher probability
- select a representative population of the higher probability particles
- include enough lower probability particles to give the filter a chance of detecting strongly nonlinear behavior.

Resampling Algorithms

*In practical applications of sequential Monte Carlo methods, residual, stratified, and systematic resampling are generally found to provide comparable results. ***Despite the lack of complete theoretical analysis of its behavior*, systematic resampling is often preferred because it is the simplest method to implement.***

Comparison of Resampling Schemes for Particle Filtering 0507025 by:
Randal Douc, Oliver Cappe, Eric Moulines.

<https://arxiv.org/pdf/cs/0507025.pdf>

Resampling Algorithm: Multinomial

```
print('cumulative sume is', np.cumsum([.1, .2, .1, .6]))  
plot_cumsum([.1, .2, .1, .6])
```

cumulative sume is [0.1 0.3 0.4 1.0]



multinomial resampling



This is an $O(n \log(n))$ algorithm. That is not terrible, but there are $O(n)$ resampling algorithms with better properties v

Resampling Algorithm: Multinomial

- The performance of the multinomial resampling could be bad.
- Very large weight can be not sampled at all
- Smallest weight can be sampled twice.
- Rarely used in the literature or for real problems

Resampling Algorithm: Residual

- Normalized weights are multiplied by N
- Integer value = how many samples of that particle will be taken.
- $w = 0.0012$, $N = 3000$, $N * w = 3.6$
- 3 samples will be taken of that particle.
- Residual = 0.6 <-- use this with other sampling methods

residual resampling



Resampling Algorithm: Residual

- Ensures all the largest weights are resampled multiple times
- Doesn't necessarily evenly distribute the samples across the particles

Resampling Algorithm: Systematic

- The space is divided into N divisions
- Choose a random offset to use for all of the divisions
- Each sample is exactly $1/N$ apart.

systematic resampling



Resampling Algorithm: Systematic

- Ensures we sample from all parts of the particle space

Resampling Algorithm: Stratified

- Make selections relatively uniformly across the particles
- Divide the cumulative sum into N equal sections
- Select one particle randomly from each section. T
- This guarantees each sample 0 and $2/N$ apart.

stratified resampling



Resampling Algorithm: Stratified

- a bit better than systematic sampling at ensuring the higher weights get resampled more.

Resampling Algorithms

std::discrete_distribution

Defined in header `<random>`

```
template< class IntType = int >      (since C++11)  
class discrete_distribution;
```

std::discrete_distribution produces random integers on the interval $[0, n)$, where the probability of each individual integer i is defined as w_i/S , that is the *weight* of the i th integer divided by the sum of all n weights.

std::discrete_distribution satisfies all requirements of [RandomNumberDistribution](#)

Resampling Algorithms

```
int main()
{
    std::random_device rd;
    std::mt19937 gen(rd());
    std::discrete_distribution<> d({40, 10, 10, 40});
    std::map<int, int> m;
    for(int n=0; n<10000; ++n) {
        ++m[d(gen)];
    }
    for(auto p : m) {
        std::cout << p.first << " generated " << p.second << " times\n";
    }
}
```

Output:

```
0 generated 4028 times
1 generated 978 times
2 generated 1012 times
3 generated 3982 times
```

Outline

- What is it?
- Real world examples
- Key Concepts
- The PF Algorithm
- PF Algorithm Review
- Resampling Algorithms
- **Compare PF with other Probabilistic Filters**

Quiz

	state space	belief	efficiency	in robotics
Class 1 Histogram Filters	<input checked="" type="checkbox"/> Discrete <input type="checkbox"/> Continuous	<input type="checkbox"/> Unimodal <input checked="" type="checkbox"/> Multimodal	<input type="checkbox"/> Quadratic <input checked="" type="checkbox"/> Exponential	<input type="checkbox"/> Exact <input checked="" type="checkbox"/> Approximate
Class 2 Kalman Filters	<input type="checkbox"/> Discrete <input checked="" type="checkbox"/> Continuous	<input checked="" type="checkbox"/> Unimodal <input type="checkbox"/> Multimodal	<input checked="" type="checkbox"/> Quadratic <input type="checkbox"/> Exponential	<input type="checkbox"/> Exact <input checked="" type="checkbox"/> Approximate
Class 3 Particle Filters	Continuous	multimodal	?	approximate

	MCL	EKF
Measurements	Raw Measurements	Landmarks
Measurement Noise	Any	Gaussian
Posterior	Particles	Gaussian
Efficiency(memory)	✓	✓✓
Efficiency(time)	✓	✓✓
Ease of Implementation	✓✓	✓
Resolution	✓	✓✓
Robustness	✓✓	x
Memory & Resolution Control	Yes	No
Global Localization	Yes	No
State Space	Multimodel Discrete	Unimodal Continuous

QUIZ QUESTION

How is MCL different from other localization algorithms? Select all that apply:

- ☐ MCL uses grids to localize the robot pose
- ☒ MCL uses particles to localize the robot pose
- ☐ MCL uses histograms to localize the robot pose
- ☐ MCL is restricted by a Gaussian state space distribution
- ☒ MCL can approximate almost any state space distribution

- EKF's cannot incorporate negative information
- They cannot use the fact that a robot failed to see a feature even when expected.

- **Sebastian Thrun**

Outline

- What is it?
- Real world examples
- Key Concepts
- The PF Algorithm
- KLD Sampling (Preview)
- PF Algorithm Review
- Resampling Algorithms

Sources 1

- UPenn - <http://fltenth.org/session3#l1-3>
- MIT RACE CAR
https://www.youtube.com/watch?v=-c_0hSjgLYw
- RLABBE RESAMPLING
<https://github.com/rlabbe/Kalman-and-Bayesian-Filters-in-Python/blob/master/12-Particle-Filters.ipynb>
- Particle Filter Explained without Equations
<https://www.youtube.com/watch?v=aUkBalzMkv4>

Sources 2

- Comparison of Resampling Schemes for Particle Filtering 0507025*
by: Randal Douc, Oliver Cappe, Eric Moulines.
<https://arxiv.org/pdf/cs/0507025.pdf>
- Particle Filters in Robotics In Proceedings of Uncertainty in AI (UAI)
2002 Sebastian Thrun
- Udacity Robotics Software Engineer Nanodegree
- Udacity Intro to AI for Robotics Course
- <https://github.com/mithi/particle-filter-prototype>